


Encoder (BERT-style)

Bidirectional: every token attends to the whole sentence → understanding, classification, features

Decoder (GPT-style)

Causal mask: token t attends only to tokens $\leq t$ → autoregressive generation, summarisation