


The LLM inference ecosystem for financial NLP

From research workflow to silicon: four layers and two cross-cutting cost levers



Read top-to-bottom: the application calls an API endpoint, which is served by a runtime, which runs on silicon. Read bottom-to-top: capability flows up the stack. The two horizontal boxes inside Layer 1 (Batch API, Prompt caching) are the dominant cost levers in production financial NLP workflows; they cut across all providers and runtimes.

FIGURE 23.1: A four-layer view of the LLM inference ecosystem